

Haomin (Harmin) Qi

harminchee.github.io

INTERESTS

Open to roles in **Applied/Research Scientist, Machine Learning Engineer, AI/Agent Developer, Software Engineer**, and broader **IT roles** (technical support, platform, data). Technical interests: coding agents and multi-agent planning, LLM inference and training, multimodal, LLM security and safety-aware code generation.

EDUCATION

University of California San Diego M.S. in Electrical and Computer Engineering GPA: 3.8 /4.0 Track in Machine Learning and Data Science	Sep. 2025 – Nov. 2026 La Jolla, CA
The Chinese University of Hong Kong B.S. in Mathematics and Information Engineering Double Major Graduation Elite Stream Class	Sep. 2021 – Jul. 2025 Hong Kong SAR
University of Leeds Abroad Exchange in Computer Science GPA: 4.00 /4.00 First Class Honour	Jan. 2024 – Jul. 2024 Leeds, UK

EXPERIENCE

Shang Data Lab, UC San Diego Research Assistant Supervisor: Jingbo Shang - Developed PoCGYM, a benchmark construction framework for realistic bug injection into production Java repositories. Grounded injection in real execution traces and bug-pattern specifications, with optional semantic transform stages (call-stack deepening, aliasing, predicate reframing) to control artifact difficulty and realism - Built an end-to-end pipeline spanning injection, Maven/Surefire validation, and archival; defined two agent tasks—Trigger (write an exposing test) and Detect (localize bug from a pattern description)—on real-code bugs	Sep. 2025 – Present La Jolla, CA
Advanced Wireless Systems Group, CUHK Research Assistant Supervisor: Soung Chang Liew - Led VeriRAG, the first RAG framework for RTL DFT compliance repair; trained a contrastive autoencoder for structure-aware Verilog retrieval with an iterative Xcelium+LEC revision loop, achieving 7.72× repair rate over zero-shot GPT-o1 - Designed TopoEdge, an edge agentic framework for SDN configuration generation/repair; trained a contrastive GCN over 1,879 FRRouting topology graphs, grounding a three-agent generate–verify–repair loop on a Pi cluster at Pass@20=0.890	Apr. 2024 – Sep. 2025 Hong Kong SAR
Deloitte Machine Learning Developer Intern - Built ABA-RAG, fusing multimodal IoT signals (BVP, GSR, temperature, acceleration) via a Transformer classifier (73% acc., 0.90 recall) with a three-layer RAG pipeline using LoRA-GA fine-tuned Llama-3.1-70B, scoring 9.3/10 among four frontier LLMs - Deployed as a web platform with task generation, therapy management, and IoT dashboards; pilot across $N=10$ learners yielded expert ratings of 9.63/10 (ICC=0.76), on par with practitioner-led interventions	Sep. 2024 – Dec. 2024 Hong Kong SAR
Wireless Ad-Hoc & Sensor Networks Lab, NCU Research Intern Supervisor: Min-Te Sun - Proposed Hybrid PEFT, fusing LoRA-GA and BOFT per layer via gradient-norm-based adaptive mixing; reduced training time by 2.1× and peak GPU memory by 50% across 7B–405B models, outperforming individual baselines on GLUE, GSM8K, MT-Bench, and HumanEval - Pioneered uRNN-adapted transformer fine-tuning by embedding structured unitary matrices into attention and FFN layers with manifold-preserving gradient updates; achieved consistent gains over LoRA and BOFT on XNLI and FLORES	Jun. 2024 – Sep. 2024 Taoyuan, TW
Intell-Pro Global R-Guardian Co-Founder & AI Lead - Architected a trademark similarity engine combining CNN-based image feature extraction, template matching, and reverse image search over scalable cloud infrastructure for real-time retrieval across multiple national IP databases - Co-founded Intell-Pro Global Limited, serving law firms and IP agencies across the US, Europe, and Asia; secured HK\$675,000 in TSSSU funding and the HK Tech300 Entrepreneurship Award	May. 2023 – Present Hong Kong SAR

- Contributed to **DeepMAD**, a training-free NAS framework formulating CNN design as entropy maximization subject to effectiveness constraints; validated on ImageNet, achieving **82.8% top-1 accuracy** at **50% fewer FLOPs** vs. conventional baselines

PUBLICATIONS

PoCGYM: Evaluating AI Coding Agents for PoC Generation via Verifiable Bug Injection

Haomin Qi, Xiangzhe Xu, Yiming Huang, Jingbo Shang, Chengpeng Wang

Under review at Neural Information Processing Systems (NeurIPS 26')

Topology-Grounded Agentic Framework for Edge Networking Code Generation and Repair

Haomin Qi, Bohan Liu, Zihan Dai, Yunkai Gao, Soung Chang Liew

IEEE Wireless Communications and Networking Conference (IEEE WCNC 26')

VeriRAG: A Retrieval-Augmented Framework for Automated RTL Testability Repair

Haomin Qi*, Yuyang Du*, Lihao Zhang, Soung Chang Liew, Kexin Chen, Yining Du

International Symposium on Quality Electronic Design (ISQED 26')

GraphCue for SDN Configuration Code Synthesis

Haomin Qi, Fengfei Yu, Chengbo Huang

IEEE Consumer Communications & Networking Conference 2026 (IEEE CCNC 26')

Governance-Aware Hybrid Fine-Tuning for Multilingual Large Language Models

Haomin Qi, Chengbo Huang, Zihan Dai, Yunkai Gao, Min-Te Sun

IEEE International Conference on Big Data 2025 Workshop LLM4All (IEEE BigData 25'LLM4All)

Hybrid and Unitary PEFT for Resource-Efficient Large Language Models

Haomin Qi, Zihan Dai, Chengbo Huang

American Journal of Computer Science and Technology (AJCST)

Transforming ABA Therapy: An IoT-Guided, Retrieval-Augmented LLM Framework

Haomin Qi, Chung-Ho Sin, Rosanna Yuen-Yan Chan, Victor Chun-Man Wong

IEEE Access

PROJECTS

DAP: Panoramic Metric Depth Foundation Model

- Built a panoramic metric depth foundation model on a **2M-sample data engine** (UE5 AirSim360, Structured3D, 1.7M real panoramas) via a three-stage semi-supervised pipeline with Scene-Invariant and Realism-Invariant labelers bridging synthetic/real domain gaps
- Proposed a DINOv3-Large backbone with plug-and-play range mask heads and distortion-aware multi-loss optimization; achieved SOTA zero-shot metric depth on Stanford2D3D, Matterport3D, and Deep360

MWRCNN: Multi-Stage Wavelet Dynamic CNN for Image Restoration

- Extended MWDCNN to a unified restoration framework (deblurring, JPEG artifact removal, super-resolution, low-light, mixed degradation) by redesigning the Dynamic Convolutional Block's Weight Generator to preserve spatial context beyond global average pooling
- Outperformed TV and MWDCNN on BSD68: **>3.5 dB PSNR** at $\sigma=50$, **>0.5 dB** on GoPro real deblurring, and **50% LPIPS reduction** under mixed degradation

Adaptive Test-Time Scaling for LLM Safety Guards

- Studied parallel, sequential, and latent test-time scaling for GuardReasoner on WildJailbreak; led parallel scaling via DeepConf confidence-weighted aggregation, raising F1 to **0.759** at $k=8$, outperforming majority voting at all sample budgets
- Designed confidence-triggered adaptive policies; a two-level policy at $\tau=0.8$ achieved **92% accuracy at 63% of peak compute** vs. 91% static full-budget baseline; diffusion latent scaling reached **91% accuracy** at 50 inference steps

SKILLS

Languages: Python, Java, C, SQL, Bash, HTML, JavaScript/TypeScript, Verilog, P4

LLM Training: PyTorch, Hugging Face Transformers, PEFT, LoRA/LoRA-GA/BOFT, Accelerate, DeepSpeed, vLLM

Agents: LangGraph, LangChain, Claude Code, Codex, OpenAI API, Anthropic API, AWS Bedrock, RAG

Tools & Infra: Maven/Surefire, JUnit, Pytest, Docker, Git, Linux, FastAPI, MLflow, Weights & Biases, CI/CD